

COLUMBIA UNIVERSITY EEME E6911 FALL '25

TOPICS IN CONTROL : PROBABILISTIC ROBOTICS

MONTE CARLO LOCALIZATION (MCL)

Instructor: Ilija Hadzic

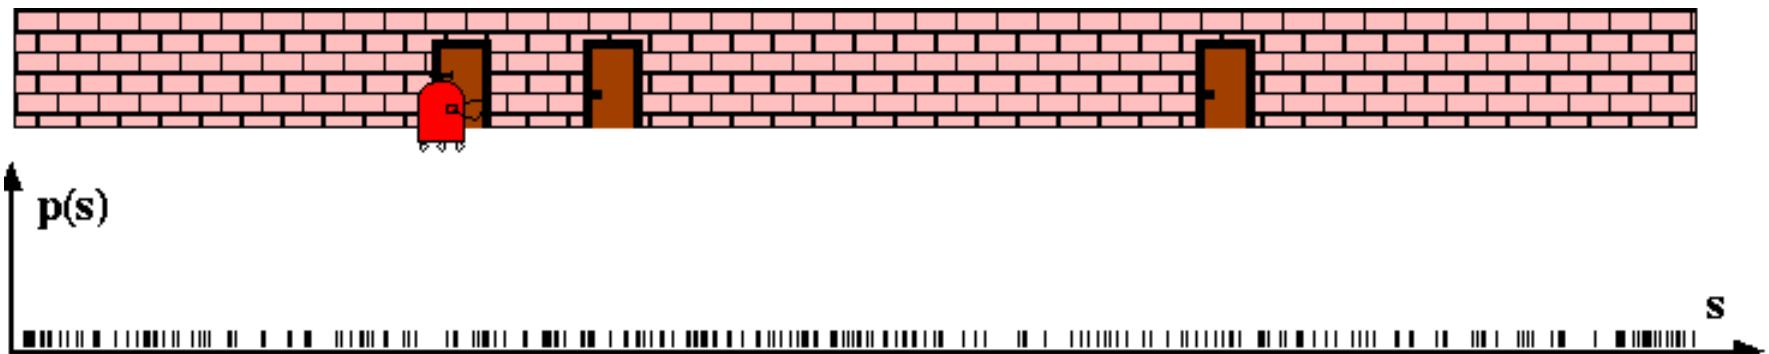
What is MCL?

- Localization using Particle Filter
- Motion Model:
 - Moves particles according to robot kinematics
 - Adds randomness according to the error model
- Measurement Model:
 - Scores particles using landmarks or raw measurement
 - Resamples according to importance
- Performance Functions:
 - Decide how many particles to use
 - Decide when to reseed the particle

Why MCL?

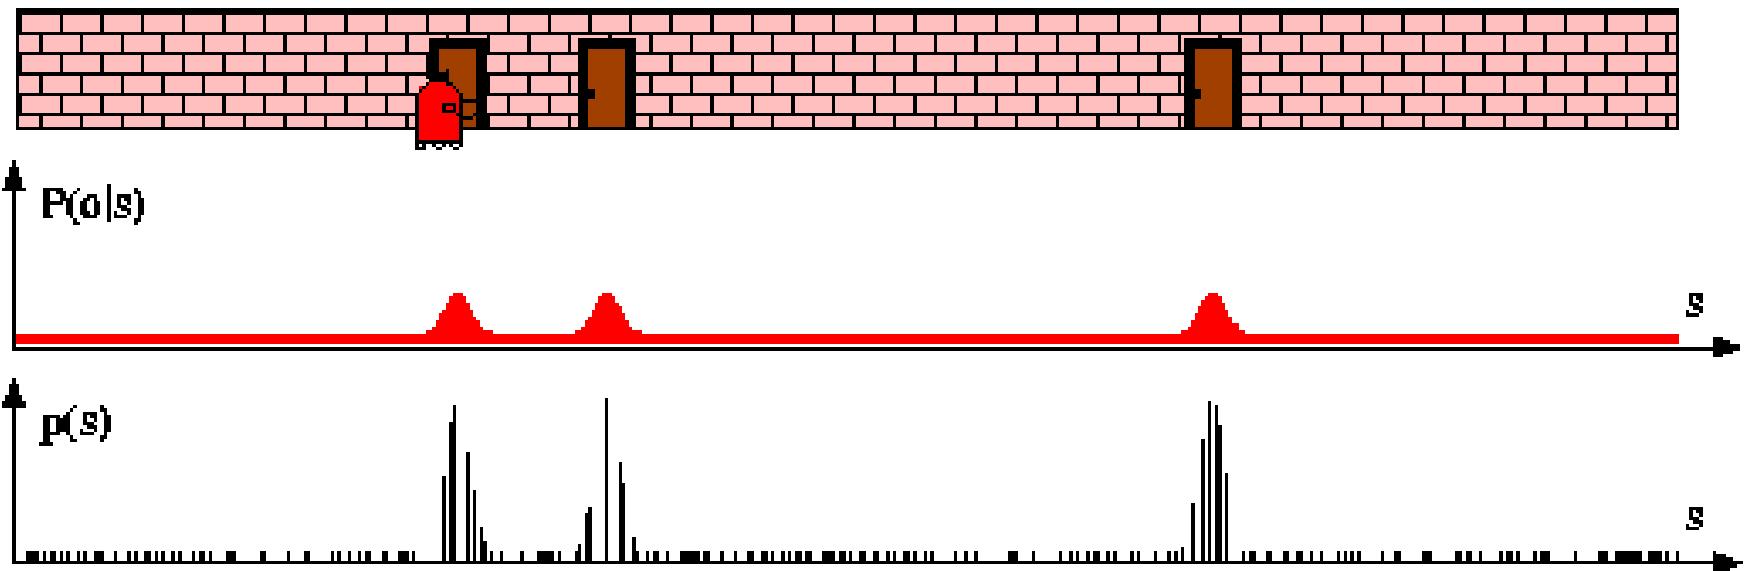
- Advantages over EKF:
 - Handles any error model
 - Works for multi-modal distributions
 - Can solve kidnap
 - Can do global localization
- Disadvantages:
 - Computationally complex (needs many particles)
 - Missing the deadline vs. degrading the performance
 - A few surprising effects (to be discussed)

1D Example From the Textbook



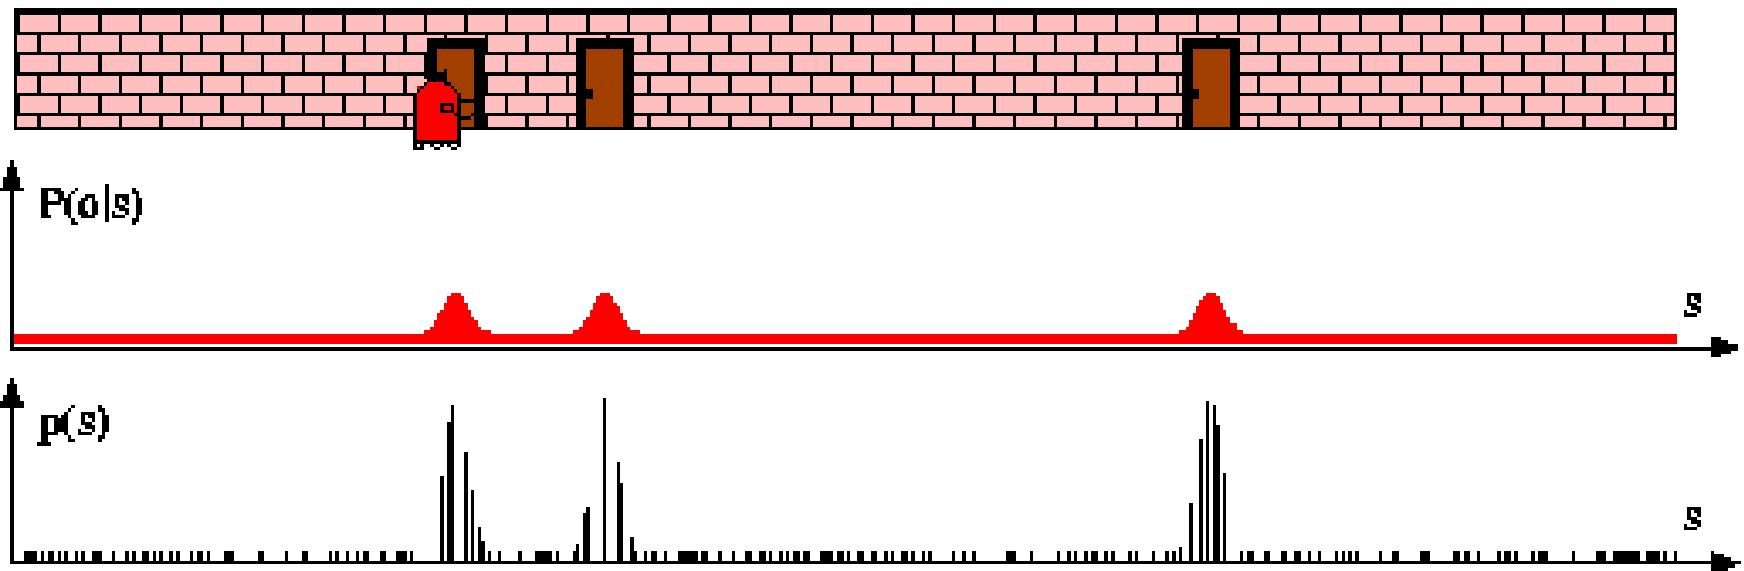
- Particles scattered at random (uniform distribution)
- Robot has no idea where it is

1D Example From the Textbook



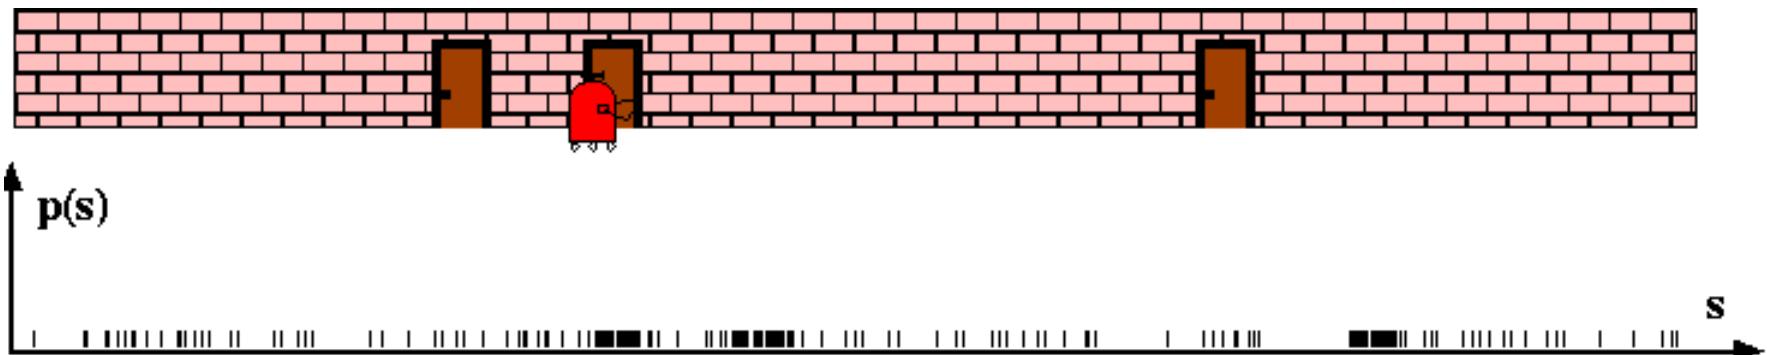
- Particles scored after observing the door
- Some particles are more important than others
- Unless the robot moves, it cannot tell anything more

1D Example From the Textbook



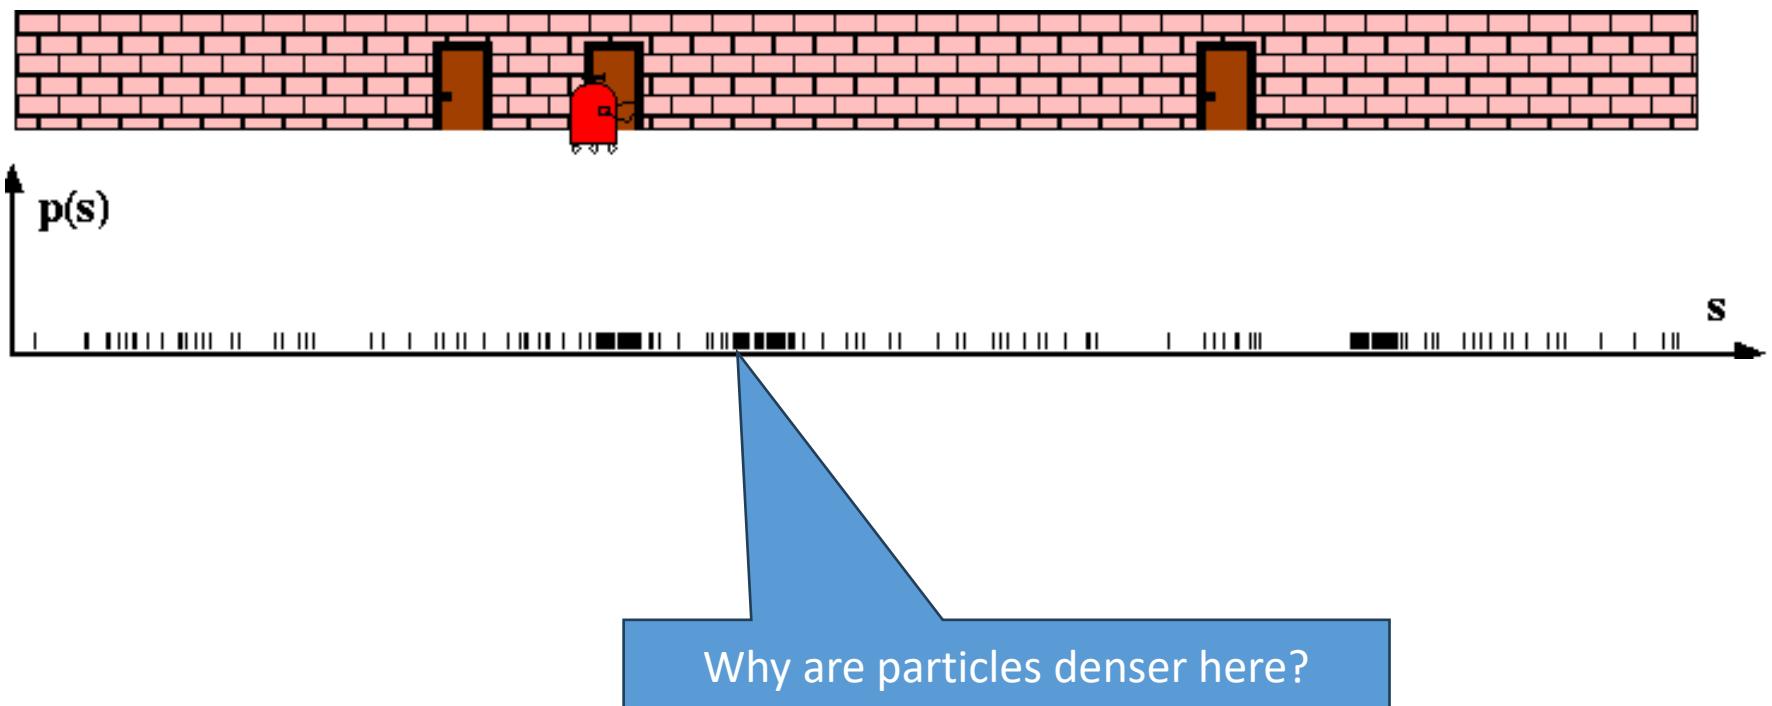
- What would the scoring look like if robot were to observe the wall instead of the door?

1D Example From the Textbook

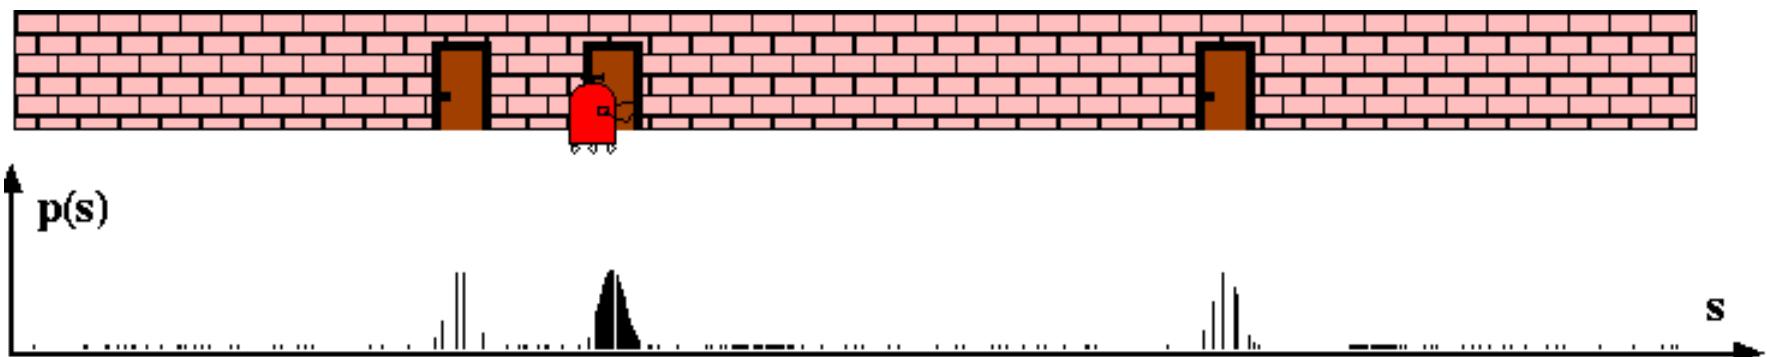


- Particles resampled
- Robot moved; particles scored again
- Resampled
- Moved again, encountered the second door

1D Example From the Textbook

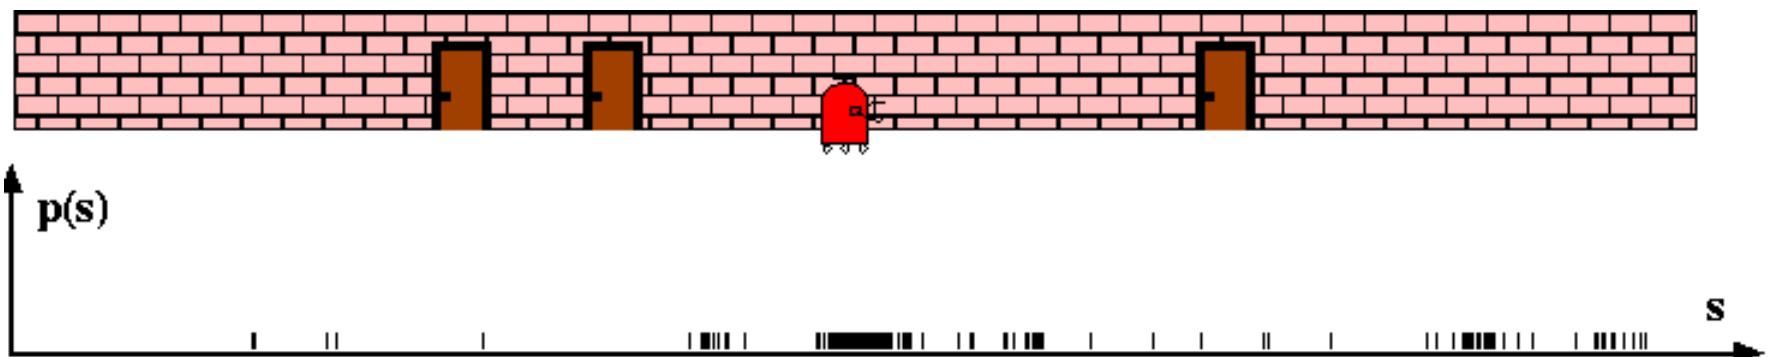


1D Example From the Textbook



- Particles scored
- Resampled
- Robot moves again

1D Example From the Textbook

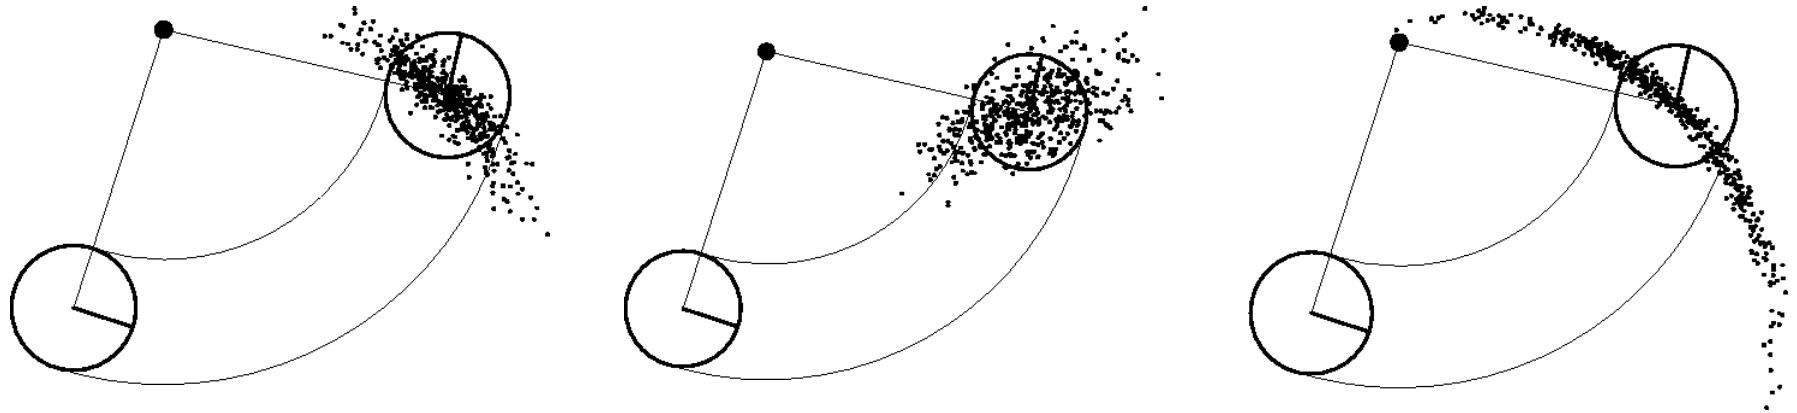


- Distribution is now (almost) unimodal
- All particles bunched around robot's true position
- Other particles died off

2D Motion Model Example

```
1:   Algorithm sample_motion_model_velocity( $u_t, x_{t-1}$ ):  
2:      $\hat{v} = v + \text{sample}(\alpha_1|v| + \alpha_2|\omega|)$   
3:      $\hat{\omega} = \omega + \text{sample}(\alpha_3|v| + \alpha_4|\omega|)$   
4:      $\hat{\gamma} = \text{sample}(\alpha_5|v| + \alpha_6|\omega|)$   
5:      $x' = x - \frac{\hat{v}}{\hat{\omega}} \sin \theta + \frac{\hat{v}}{\hat{\omega}} \sin(\theta + \hat{\omega}\Delta t)$   
6:      $y' = y + \frac{\hat{v}}{\hat{\omega}} \cos \theta - \frac{\hat{v}}{\hat{\omega}} \cos(\theta + \hat{\omega}\Delta t)$   
7:      $\theta' = \theta + \hat{\omega}\Delta t + \hat{\gamma}\Delta t$   
8:     return  $x_t = (x', y', \theta')^T$ 
```

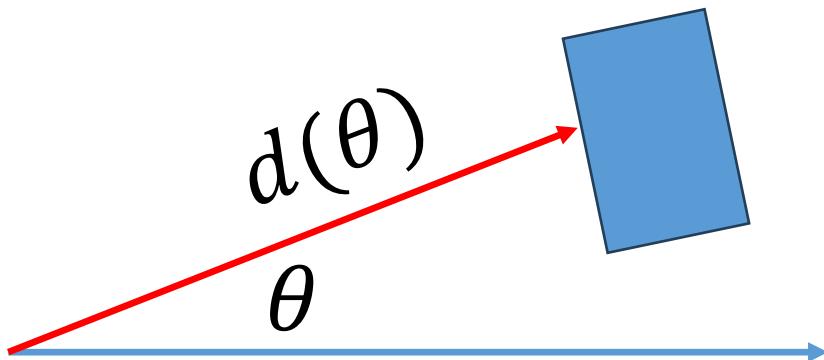
2D Motion Model Example



- Thrun's model we saw before, but this time using particles
- Only looks at velocity command

Range Finder Measurement Model

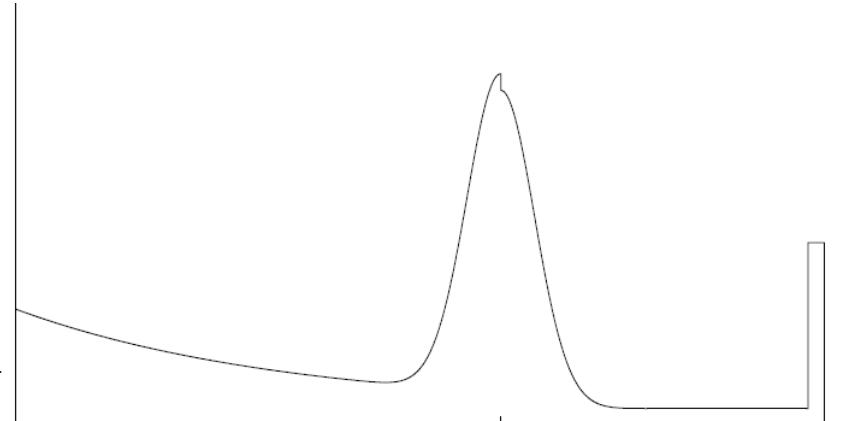
- Senses the distance to a fixed object (e.g. walls):
 - LIDAR, Radar, Sonar
- Output is the distance in polar coordinates



Beam Model

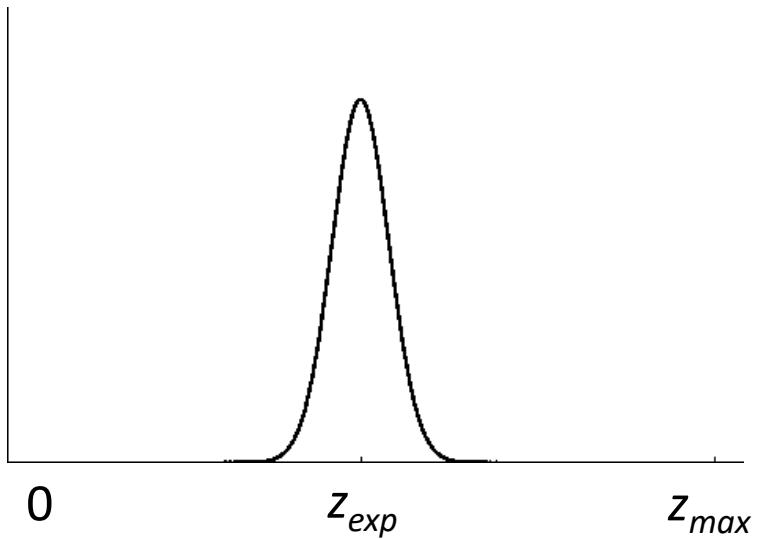
- What are we modeling?
 - $p(z_i[n] | \bar{x}[n], m)$ for each beam i .

$$P(z | x, m) = \begin{pmatrix} \alpha_{\text{hit}} \\ \alpha_{\text{unexp}} \\ \alpha_{\text{max}} \\ \alpha_{\text{rand}} \end{pmatrix}^T \cdot \begin{pmatrix} P_{\text{hit}}(z | x, m) \\ P_{\text{unexp}}(z | x, m) \\ P_{\text{max}}(z | x, m) \\ P_{\text{rand}}(z | x, m) \end{pmatrix}$$

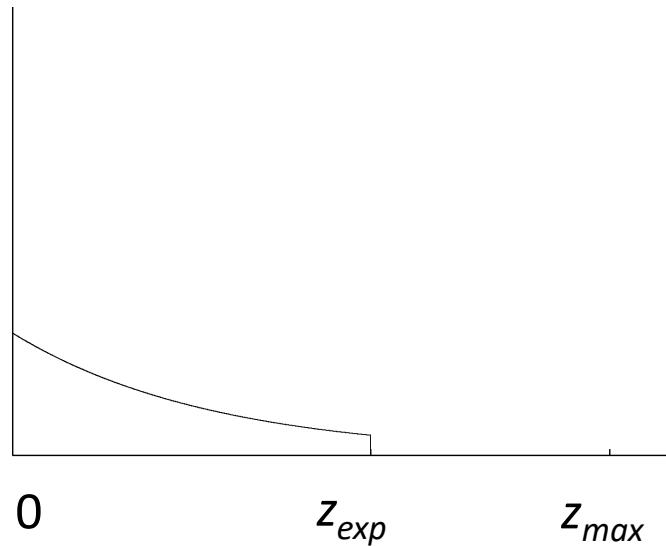


Beam Model

Measurement noise



Unexpected obstacles

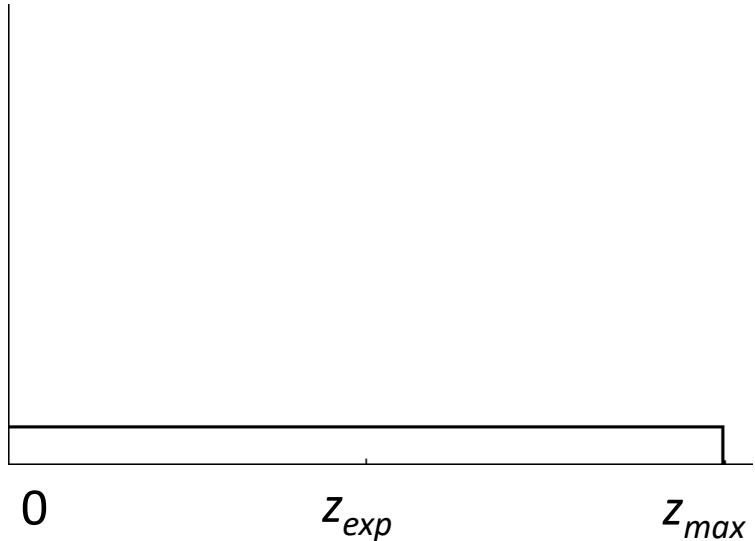


$$P_{hit}(z | x, m) = \eta \frac{1}{\sqrt{2\pi b}} e^{-\frac{1}{2} \frac{(z - z_{exp})^2}{b}}$$

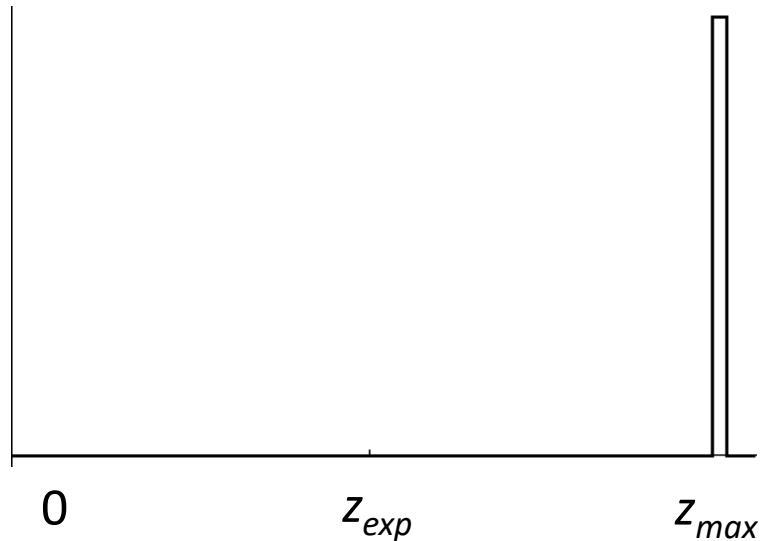
$$P_{unexp}(z | x, m) = \begin{cases} \eta \lambda e^{-\lambda z} & z < z_{exp} \\ 0 & otherwise \end{cases}$$

Beam Model

Random measurement



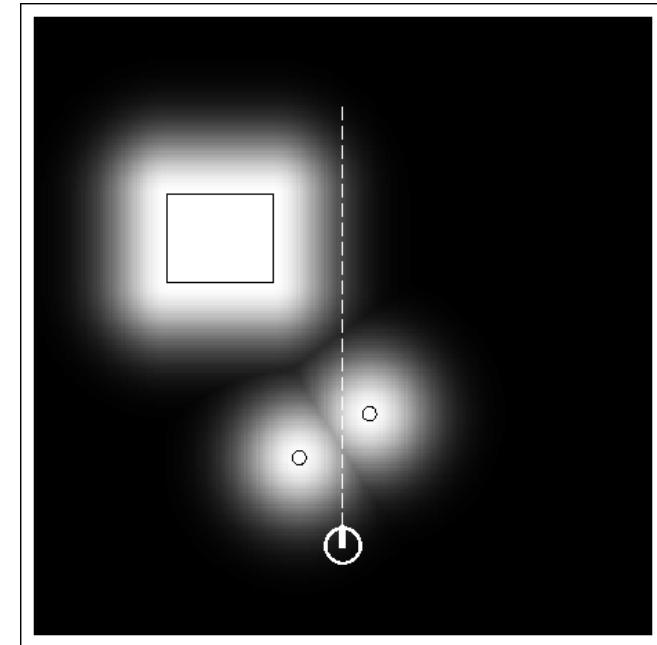
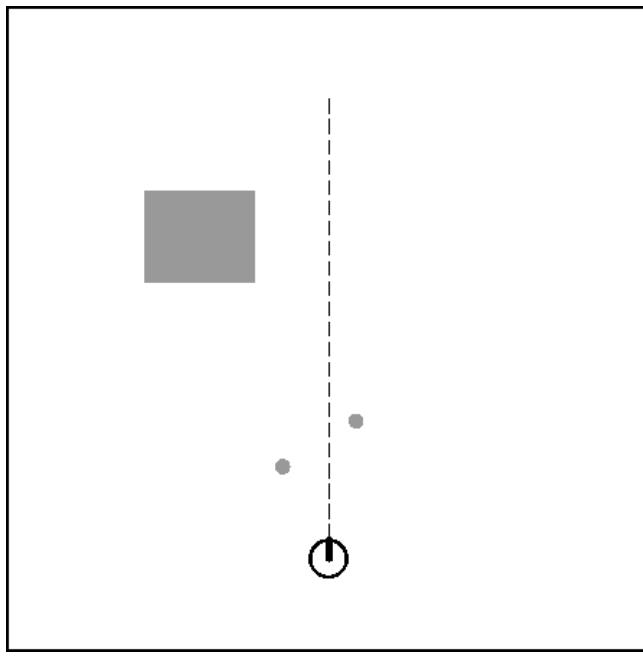
Max range



$$P_{rand}(z | x, m) = \eta \frac{1}{z_{max}}$$

$$P_{\max}(z | x, m) = \eta \frac{1}{z_{small}}$$

Likelihood Fields



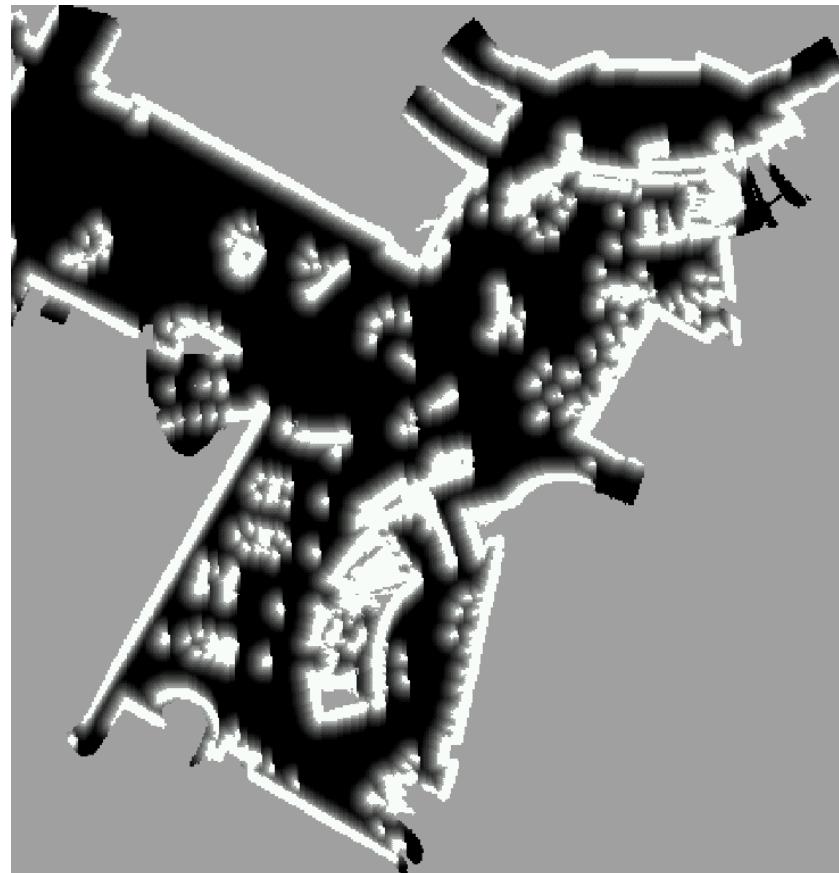
Likelihood field



San Jose Tech Museum



Occupancy grid map

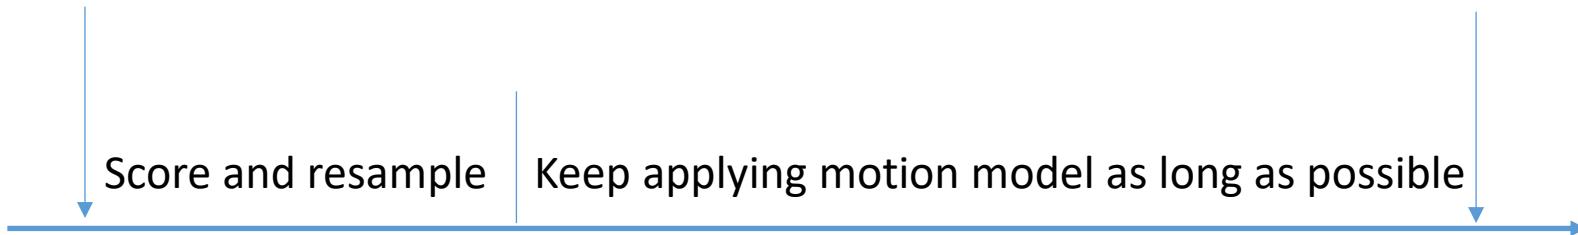


Likelihood field

Sampling Against the Deadline

Measurement came in.
Stop producing samples
Throw away unused samples

Measurement came in.
Stop producing samples
Throw away unused samples

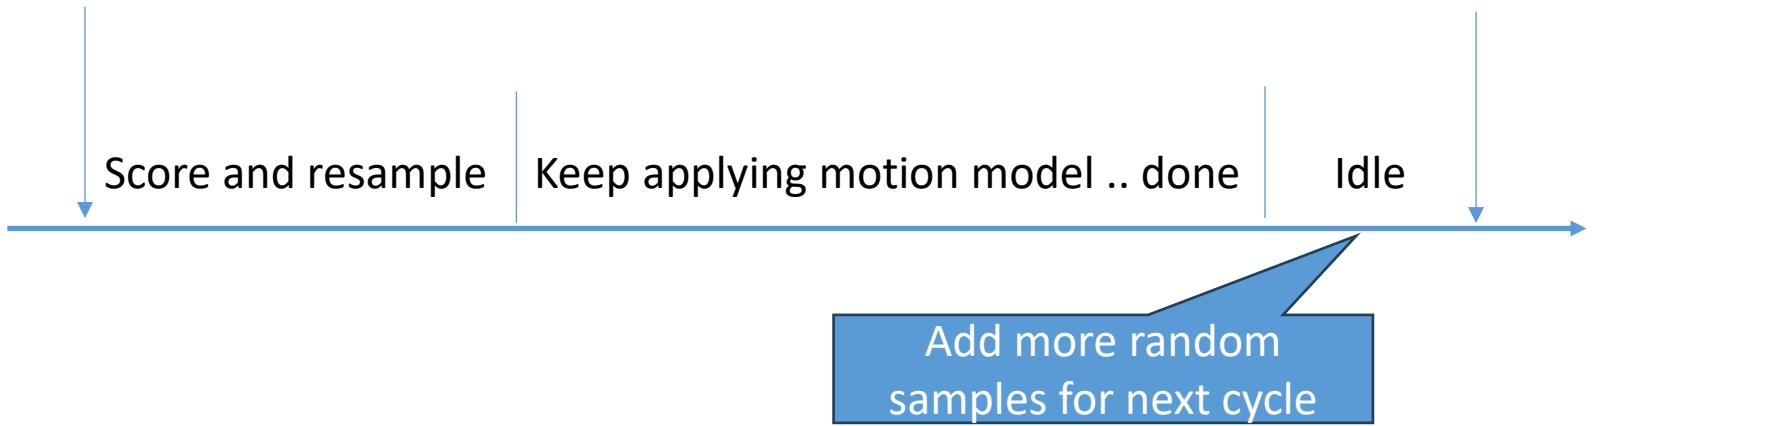


- Keep sampling until the new measurements comes in.
- Work with particles produced thus far
- Faster CPU achieves better precision.

Sampling Against the Deadline

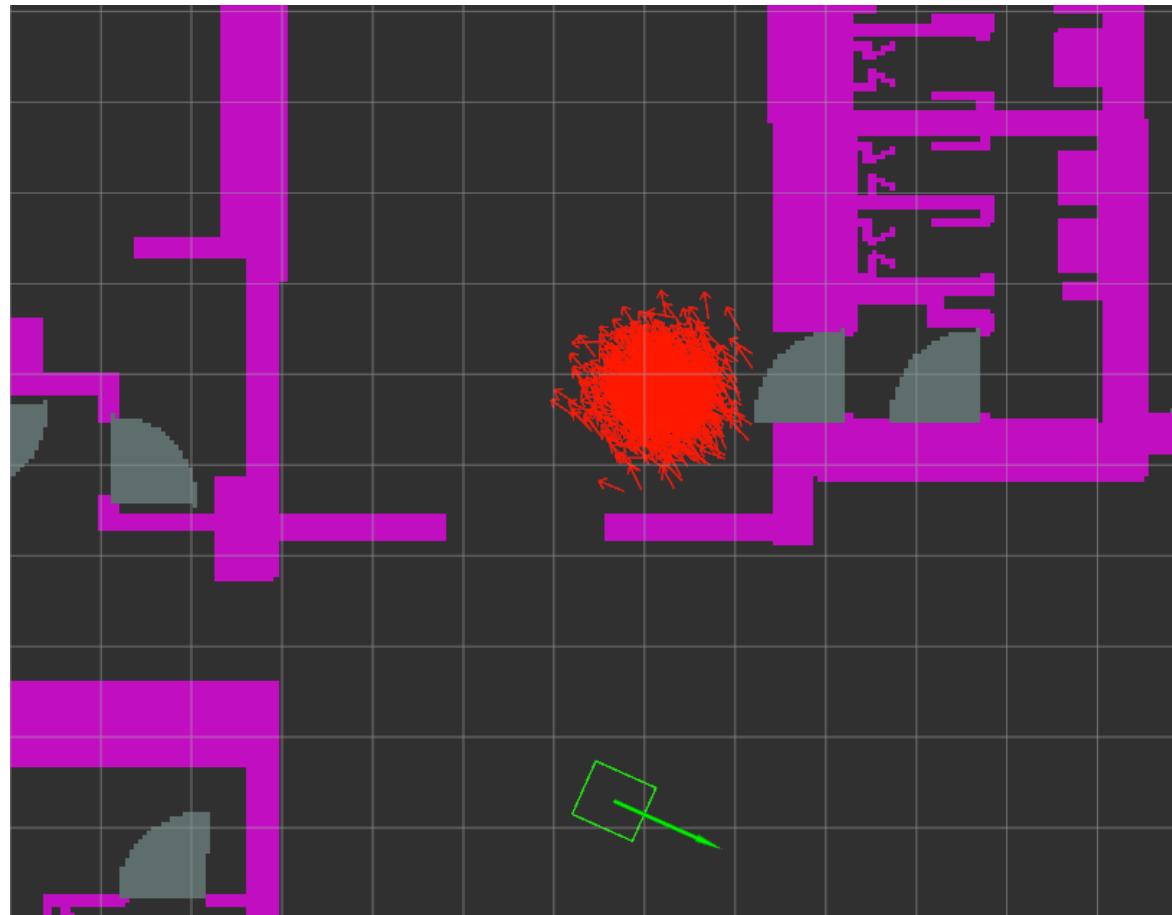
Measurement came in.
Stop producing samples
Throw away unused samples

Measurement came in.
Stop producing samples
Throw away unused samples

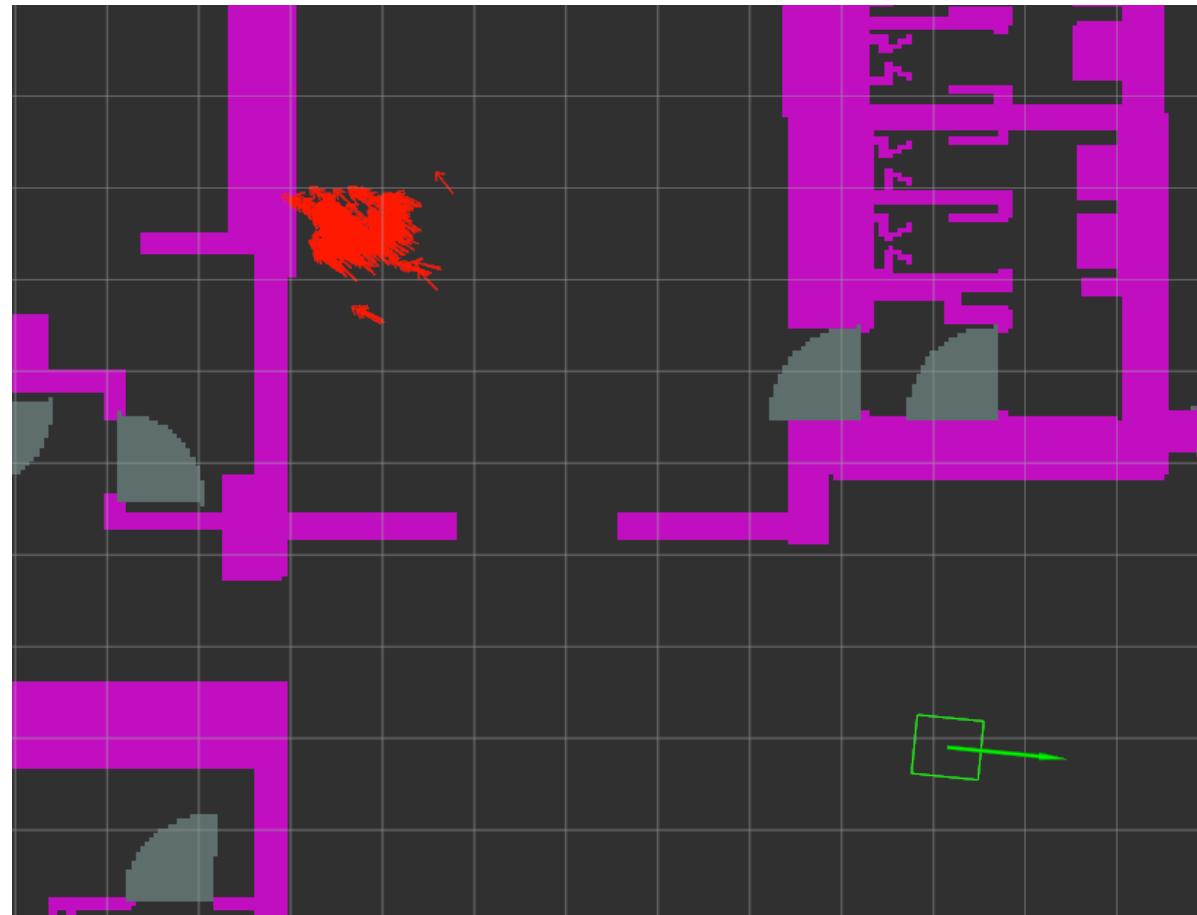


- If we have time budget left, add more samples.
- The number of samples converges to fill the available CPU time.
- Algorithm adapts to available compute resources.

Adding Random Particles



Adding Random Particles



Adding Random Particles

- Problem:
 - Estimate remains stuck to the initial cloud
 - No way to get out
- Solution:
 - Add random particles
 - How many?
 - Where?

Adding Random Particles

$$w^{[j]}[n] = p(z[n] | \bar{x}^{[m]}[n])$$

$$w_{\text{avg}}[n] = \frac{1}{N} \sum_{j=1}^N w^{[j]}[n]$$

$$w_{\text{slow}} = w_{\text{slow}} + \alpha_{\text{slow}}(w_{\text{avg}} - w_{\text{slow}})$$

$$w_{\text{fast}} = w_{\text{fast}} + \alpha_{\text{fast}}(w_{\text{avg}} - w_{\text{fast}})$$

$$0 < \alpha_{\text{slow}} \ll \alpha_{\text{fast}}$$

Adding Random Particles

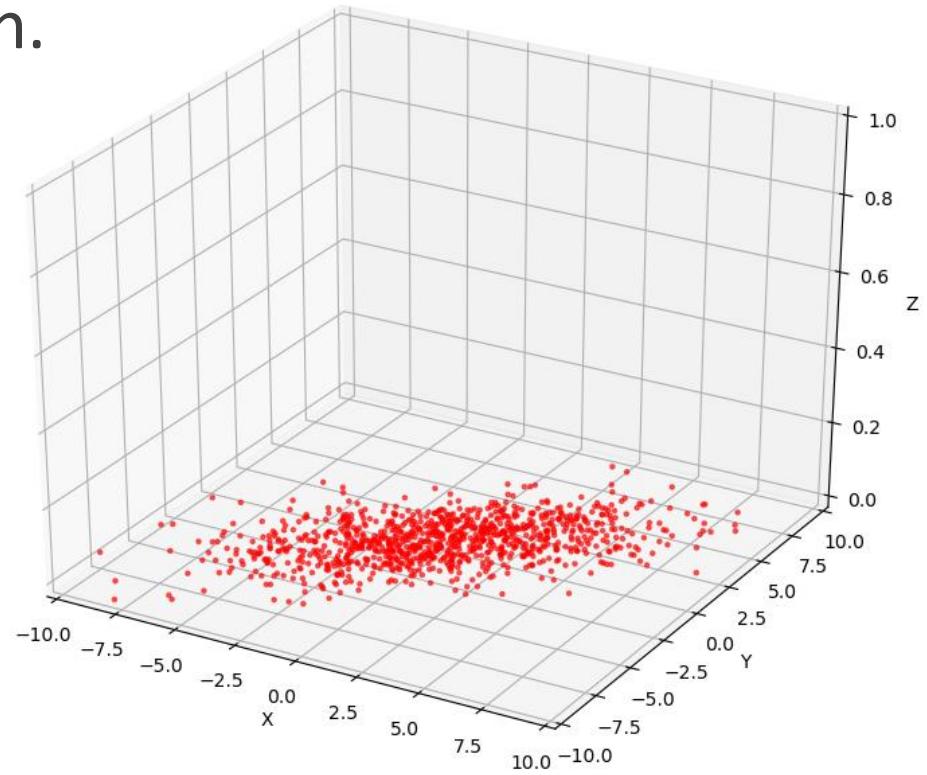
- Probability of replacing a traced particle with the random particle

$$\max \left\{ 0, 1 - \frac{w_{\text{fast}}}{w_{\text{slow}}} \right\}$$

- Picking the area to cover with random particles is usually based on heuristics.
- Landmark identification can help:
 - Example: Spotting the Statue of Liberty narrows down the area to a few streets on Manhattan or in New Jersey

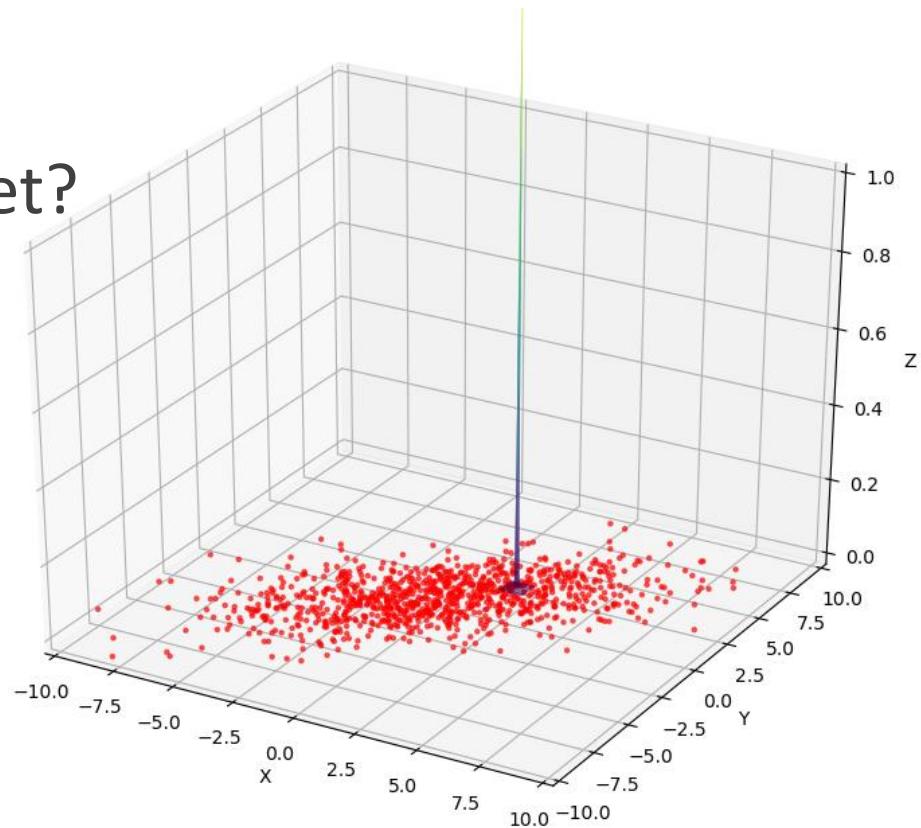
Perfect Sensor Paradox

- Consider a particle cloud.
- Measurement comes in.
- Score them!



Perfect Sensor Paradox

- Perfect measurement!
- What is the variance?
- What weights do we get?
- Why?



Remedies

- Add artificial noise to the sensor.
- Explicitly recognize the case and defend from it:
 - Take the measurement at the face value
 - Initialize the particle cloud
- Mid-ground solution:
 - For a subset of particles, use measurement distribution
 - Do the regular processing of the rest
 - Problem: This can be computationally expensive
- Exercise:
 - Take the code from previous class
 - Keep lowering the measurement variance until it crashes
 - Explain the crash and fix it!

KLD Sampling

- Kullback-Leibler Divergence:
 - Measure of how different two distributions are.
- In theory:
 - Measures sampled and true distributions
 - If the difference is high, add more particles.
- In practice:
 - Watch where particles are landing.
 - If they are landing at empty spots, keep adding.
 - If they are landing at similar spots, stop.

KLD Sampling

- Parameters: δ, ϵ
- Set some minimum number of samples M_{\min}
- Partition the space into bins
- Count the number of bins with at least one sample
- Calculate this:

$$M_x = \frac{k-1}{2\epsilon} \left(1 - \frac{2}{9(k-1)} + \sqrt{\frac{2}{9(k-1)}} z_{1-\delta} \right)^3$$

- We need $\max(M_x, M_{\min})$ samples.

ROS: AMCL Node

min_particles	0	1000	500	laser_min_range	-1.0	1000.0	-1.0
max_particles	0	10000	5000	laser_max_range	-1.0	1000.0	50.0
kld_err	0.0	1.0	0.05	laser_max_beams	0	250	180
kld_z	0.0	1.0	0.99	laser_z_hit	0.0	1.0	0.5
update_min_d	0.0	5.0	0.2	laser_z_short	0.0	1.0	0.1
update_min_a	0.0	6.283185307	0.2	laser_z_max	0.0	1.0	0.05
resample_interval	0	20	1	laser_z_rand	0.0	1.0	0.5
transform_tolerance	0.0	2.0	0.2	laser_sigma_hit	0.0	10.0	0.2
recovery_alpha_slow	0.0	0.5	0.0	laser_lambda_short	0.0	10.0	0.1
recovery_alpha_fast	0.0	1.0	0.0	laser_likelihood_max_dist	0.0	20.0	2.0
do_beamskip	<input type="checkbox"/>			laser_model_type	likelihood_field_const (likelihood_field)		
beam_skip_distance	0.0	2.0	0.5	odom_model_type	diff_corrected_const (diff-corrected)		
beam_skip_threshold	0.0	1.0	0.3	odom_alpha1	0.0	10.0	0.02
tf_broadcast	<input type="checkbox"/>			odom_alpha2	0.0	10.0	0.01
force_update_after_initialpose	<input type="checkbox"/>			odom_alpha3	0.0	10.0	0.01
force_update_after_set_map	<input type="checkbox"/>			odom_alpha4	0.0	10.0	0.02
gui_publish_rate	-1.0	100.0	10.0	odom_alpha5	0.0	10.0	0.02
save_pose_rate	-1.0	10.0	-1.0	odom_frame_id	odom		
use_map_topic	<input type="checkbox"/>			base_frame_id	base_link		
first_map_only	<input type="checkbox"/>			global_frame_id	map		
				restore_defaults	<input type="checkbox"/>		