


COLUMBIA UNIVERSITY EEME E6911 FALL '25

TOPICS IN CONTROL : PROBABILISTIC ROBOTICS

MARKOV DECISION PROCESS

Instructor: Ilija Hadzic

Fully Observable MDP



- Present state is known:
 - No uncertainty
- Transitions are uncertain:
 - Given the input
 - Randomness in transition
- Example:
 - Ground robot skidding
 - Drone disturbed by the wind
- Fully observable
- Partially controllable

Reward, Value, Horizon, Discount

- $r(x, u)$: reward/penalty of moving
 - Typically, small negative for just moving around.
- $V(x)$: value of the state
 - Typically, high positive for goal state(s).
 - Typically, high negative for lethal states.
 - Calculated for transitional states.
- T : planning horizon
 - How far into the future do we calculate the reward.
- γ : discount
 - Decay factor for future rewards.

Value Iteration

- Main idea:
 - Action rewards → some states more valuable than others
 - Infinite planning horizon → state values converge
 - Certain state → pick action with highest value gain
 - Uncertain action → use *expected* value gain
- Application:
 - Use as global planner
 - One time state value calculation per goal: high computation
 - Policy (action) extraction: low computation
 - Approximations for partially-observable MDPs:
 - QMDP, AMDP, MC-MDP

Value Iteration


- Assign value to each state
 - Initialize: $V_0(x) = 0, \forall x$ (usually zero, but can be anything)
 - Set the iterative update:
 - $V_{n+1}(x) = \gamma \max_u (r(x, u) + \int V_n(x^+) p(x^+ | u, x) dx^+)$
 - Let $n \rightarrow \infty$:
 - $V(x) = \gamma \max_u (r(x, u) + \int V(x^+) p(x^+ | u, x) dx^+)$
 - This is known as Bellman equation
- Pick the motion (policy extraction):
 - Maximize the *expected* value of the next state
 - $u(x) = \operatorname{argmax}_u (r(x, u) + \int V(x^+) p(x^+ | u, x) dx^+)$

Example: Quest on a Grid

Gold	Mud		
		Mud	
			Home


- Start: Home
- Goal: Find Gold
- Constraints:
 - Avoid stepping in mud
- Reward for gold: +50
- Penalty for mud: -100
- Cost of moving: -1
- Discount factor: 0.9

Motion probabilities – center square




- Requested direction:
 - $p=0.7$
- Bounce:
 - $p=0.1$
- Skid:
 - $p=0.1$
- Skid:
 - $p=0.1$

Motion probabilities – edge square



- Requested direction:
 - $p=0.8$
- Skid:
 - $p=0.1$
- Skid (or bounce):
 - $p=0.1$

Motion probabilities – corner square



- Requested direction:
 - $p=0.9$
- Skid:
 - $p=0.1$

Initial score

+50	-100	0	0
0	0	-100	0
0	0	0	0
0	0	0	0

- Mud and gold:
 - Terminal states
 - No motion from them
 - Value set to +50 and -100
- All other states:
 - Initialize to 0

Iterations

+50	-100		0
0	0	-100	0
0	0	0	0
0	0	0	0

Consider this state

- If the command is “left”:
 - $0.8 \cdot -100$
 - $0.1 \cdot 0$
 - $0.1 \cdot -100$
- Add cost of motion:
 - $-90 - 1 = -91$
- Apply discount: -81.9
- Repeat for other two commands:
 - Right: -18.9
 - Down: -81.9
- Max: $-18.9 \rightarrow$ New value

Iteration 1

+50	-100	-18.9	-0.9
35.1	-18.9	-100	-9.9
-0.9	-0.9	-9.9	-0.9
-0.9	-0.9	-0.9	-0.9

- Continue the process for other states.
- Careful:
 - Do not mix up present and updated values
 - Pick only the maximum
 - Watch the sign when comparing

Iteration 2

+50	-100	-19.55	-10.62
33.32	3.13	-100	-10.63
24.21	-4.14	-10.63	-3.33
-1.71	-1.71	-2.52	-1.71

- Repeat the process until convergence

Iteration 3

+50	-100	-26.55	-11.27
37.56	1.72	-100	-13.25
22.56	13.52	-12.16	-4.04
18.56	-2.73	-3.24	-3.24

- Repeat the process until convergence

Convergence

+50	-100	-23.53	-6.43
38.57	7.37	-100	-4.22
31.21	21.92	6.16	8.7
26.32	21.49	16.3	13.09

- Converges after 29 iteration for $\epsilon = 0.001$

Policy Extraction

+50	-100	-23.53	-6.43
38.57	7.37	-100	-4.22
31.21	21.92	6.16	8.7
26.32	21.49	16.3	13.09


- Where will the robot go?
- Left? Add this up:
 - $0.7 \cdot 31.21$
 - $0.1 \cdot 7.37$
 - $0.1 \cdot 21.49$
 - $0.1 \cdot 6.16$
- Less the cost of motion
 - $25.35 - 1 = 24.35$
- Down?
 - $18.52 - 1 = 17.52$
- Up? Right?
- Left maximizes the gain!

Policy Extraction

+50	-100	-23.53	-6.43
38.57	7.37	-100	-4.22
31	1.92	6.16	8.7
26.32	21.49	16.3	13.09

- It may be tempting to just pick the highest-value neighbor
- Must factor motion uncertainty!
- Intuitive policy:
 - Steer clear of mud-pits!
 - Value Iteration has learned that!

Partially Observable MDP




- Present state is uncertain:
 - Randomness in state belief
- Transitions are uncertain:
 - Given the input
 - Randomness in transition
- Example:
 - Ground robot skidding
 - Drone disturbed by the wind
 - Probabilistic localization
- Fully observable
- Partially controllable

Partially Observable MDP

- Must consider full state space with associated PDF
- Exact solution not tractable (see Chapter 15)
- Practical solutions are approximate:
 - QMDP
 - AMDP
 - MC-MDP

Partially Observable MDP

- Must consider full state space with associated PDF
- Exact solution not tractable (see Chapter 15)
- Practical solutions are approximate:
 - QMDP
 - AMDP
 - MC-MDP



We will only
discuss this

QMDP

- Calculate state values: same as MDP

QMDP

- Consider full state space
 - Associated with belief PDF
- Calculate expected future values for all states

$$Q(x_i, u) = r(x_i, u) + \sum_{j=1}^N V(x_j)p(x_j|u, x_i)$$

- Pick control

$$u = \operatorname{argmax}_u \sum_{i=1}^N p_i Q(x_i, u)$$

MDP

- Focus on present state x
 - Known with certainty
- Pick control
 - Maximize expected future value

$$u = \operatorname{argmax}_u \left[r(x, u) + \sum_{j=1}^N V(x_j)p(x_j|u, x) \right]$$

QMDP

- Calculate state values: same as MDP

QMDP

- Consider full state space transitions
 - Associated with mean over state-space transitions (future states)
- Calculate expected future values for all states

$$Q(x_i, u) = r(x_i, u) + \sum_{j=1}^N V(x_j)p(x_j|u, x_i)$$

- Pick control

$$u = \operatorname{argmax}_u \sum_{i=1}^N p_i Q(x_i, u)$$

Mean over state belief (present state)

MDP

- Focus on present state x
 - Known with certainty
- Pick control
 - Maximize expected future value

$$u = \operatorname{argmax}_u \left[r(x, u) + \sum_{j=1}^N V(x_j)p(x_j|u, x) \right]$$

Only one present state